

RED WINE QUALITY DATASET: ANALYSIS

REMSCHIRA

ABSTRACT

The wine quality dataset ([Cortez et al. 2009](#)) contains physicochemical and sensory information for red and white variants of the Portuguese “Vinho Verde” wine. The goal of this project is to use machine learning and data analysis methods to predict accurately the quality of a red wine based on its physicochemical makeup using the red wine dataset. The quality of the wine is measured by a score that takes integer values between 0 and 10. We apply several regression and classification methods to the dataset in order to determine which algorithm provides the most accurate predictions. For each of the tested machine learning algorithms, we search the hyperparameter space to find where it performs best on our accuracy measures. As part of our accuracy tests, we include 5-fold cross validation on the training dataset. We incorporate outlier detection and removal to understand the influence of potentially anomalous data on the results. To aid the visualization of the data we use Principal Component Analysis to project the data into two- and three-dimensions. We find that the Random Forest classifier performs best on our accuracy measures.

1. OUTLINE

The outline of this write-up is as follows. In Section 2, we visualize the distribution of the data and look for correlations between the physicochemical characteristics or features of the wine. Principal Component Analysis (PCA) is applied to see how much variance is explained by each principal component of the data. Section 3 addresses our implementation of outlier detection methods. The two algorithms we study are known as “Elliptical Envelope” (`sklearn.covariance.EllipticalEnvelope`) and “Isolation Forest” (`sklearn.ensemble.IsolationForest`).¹ Section 4 provides an outline of the various machine learning algorithms we apply to the dataset and reports the results of each method under our performance measures.

2. DATA VISUALIZATION

The red wine dataset contains data for 1599 wines. Each wine is characterized by 11 different physicochemical features and is scored on a 0 - 10 integer scale. Figure 1 shows the distribution of each variable. The horizontal axis shows that the variables are on different scales, which suggests the need to shift and scale each variable. Additionally, a by-eye inspection of the distributions suggests that some of the variables are approximately normally distributed (e.g. the “density” and “pH” variables). To check this we could use one of the popular tests of normality, such as the chi-square test. However, any test of normality for our large sample will likely fail because no variable in the dataset is exactly normally distributed. Moreover, although the pH variable looks normally distributed, for example, this may be an artifact of the binning used to make the histogram.

The boxplot for each variable is shown in Figure 2. The “free sulfur dioxide” and “total sulfur dioxide” features have a much wider range than the other variables. However, this is likely due to each variable being expressed in different units. To test this, we present the box plot graph in Figure 3 for the scaled version of the training data for the wine features.

We now would like to visualize the correlation between the wine variables. We do so with a “heat map” (Figure 4) (idea from Li 2017). In the case of a set of variables that are 100% independent, all off-diagonal entries in the heatmap would be zero. That is, the correlation between any two of the variables would be zero, assuming non-zero standard deviations. This is not the case for our dataset. However, this is not surprising as one expects correlations between the characteristics of a wine.

Our dataset can be thought of as a cloud of points in 11-dimensions (there are 11 features) and each point is labeled by that particular wine’s score. Principal component analysis (PCA) (`sklearn.decomposition.PCA`) finds in which direction the variation is greatest in the cloud of data points. This direction is the first principal component of the dataset. The procedure also finds the direction, perpendicular to the first component, along which the variation is second greatest. This is the second principal component. The procedure continues for all 11 dimensions of the dataset. Through PCA, our dataset gets expressed in its principal components that are linearly uncorrelated to each other. Each component accounts for a certain amount of the variation in the dataset. This is visualized in Figure 5 (idea for plot from K Hong and Raschka 2015). One can project the PCA-transformed dataset into lower dimensions. Figures 6 and 7 show the data projected into two and three dimensions (Raschka 2015).

3. OUTLIER DETECTION

¹ This project makes use of the sklearn Python library (Pedregosa et al. 2011).

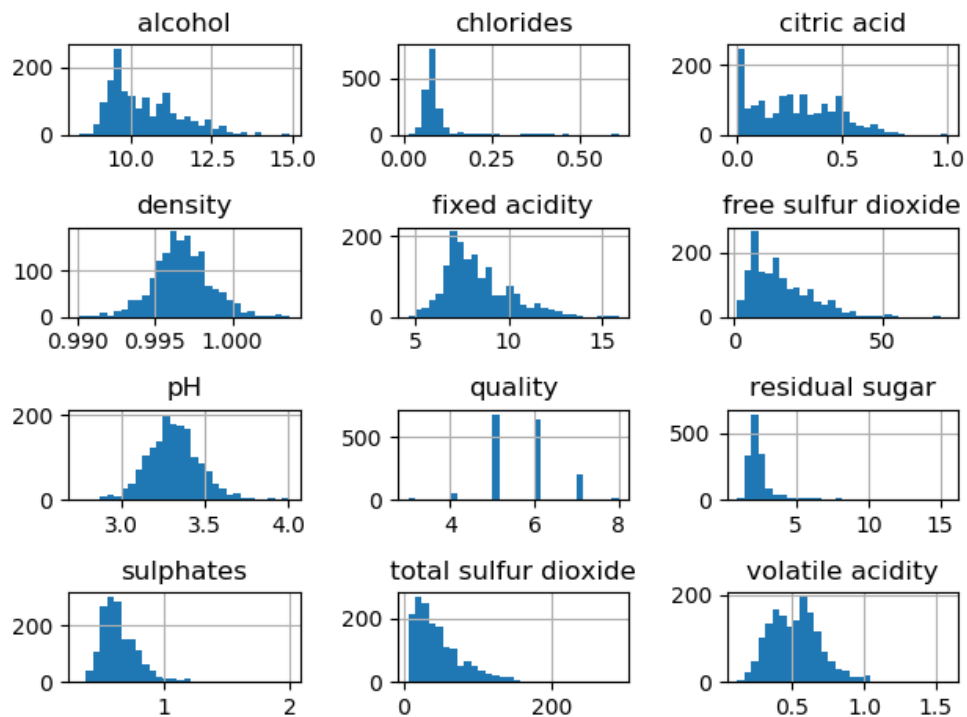


Figure 1. Binned distributions of the variables in the red wine dataset ordered alphabetically. “Quality” is the dependent variable and is categorical. All other variables are the physical attributes of the wines.




Figure 2. The unscaled distributions in boxplot-form of the variables in the red wine dataset.




Figure 3. The scaled distributions in boxplot-form of the wine features in the training set.

Here we discuss our implementation of two different outlier detection methods and discuss how the results of the two methods applied to the wines dataset differ. The first method we will address is known in the sklearn library as `sklearn.covariance.EllipticalEnvelope`. The method assumes the data to be Gaussian distributed and uses robust covariance estimation. The algorithm fits an ellipse to the data and leaves out data points that are sufficiently removed from the central mode (see the scikit learn webpage entitled “Novelty and Outlier Detection”). The method takes as a parameter the contamination fraction, which is the fraction of outliers in the dataset. Figure 8 shows the projection of the wine data on the first three principal axes with separate labels for the inliers and outliers as determined by the Elliptical Envelope method.

The second outlier detection method we implemented is known as the Isolation Forest algorithm (Liu et al. 2008). The framework for this anomaly detection method is the Random Forest algorithm. The Random Forest algorithm makes use of a multitude of decision trees to make predictions. A random forest prediction, for a classification problem, is the class that is the mode of the collection of trees and, for a regression problem, is the the mean prediction of the individual trees (see the Wikipedia page on Random forest). The insight behind the Isolation Forest algorithm is that anomalous data with distinguishable attributes are isolated in the early partitions that make up a tree. Therefore, data points with short path lengths through the tree are considered likely candidates to be outliers. The graph of the wine data projected onto the three principal axes with outliers labeled separately from the rest of the data is shown in Figure 9. By inspection, the Elliptical Envelope (Figure 8) algorithm appears to select as outliers some data point lying in the interior of the projected data cloud. Classifying these data points as outliers may only seem inappropriate due to the projections we have applied to the data to make the plot. Or, the algorithm may be failing to label




Figure 4. A heatmap visualization of the correlation matrix for the variables in the wine dataset, produced using the Seaborn Python library.

properly outliers because the assumption that the data are normally distributed is inappropriate for the wine dataset. On the other hand, the Isolation Forest (Figure 9) labels only points that are away from the data cloud.

4. APPLICATION OF MACHINE LEARNING CLASSIFIERS

We now look at how a suite of machine learning classifiers performs on the red wine dataset. To prepare the data, we first split the data into training and test sets (2/3 training, 1/3 testing). This is done randomly with `sklearn.model_selection.train_test_split`. We next apply a scaler to the data with `sklearn.preprocessing.StandardScaler`, which standardizes the data. We train each classifier on the scaled data set, and test the trained model on the test data. The performance of a classifier is measured by the accuracy (`sklearn.metrics.accuracy_score`), the confusion matrix (`sklearn.metrics.confusion_matrix`), and by the average of a 5-fold cross validation (`sklearn.model_selection.cross_val_score`). The classifiers we used for this project were: Logistic regression, Decision tree, Random forest, Adaboost, Gradient boosting, Linear SVM, and SVM with RBF kernel.

The performance of the classifiers can be found in `classifier_performance.dat`. For each model, we vary one of its hyperparameters. The best performing classifier is the Random forest with

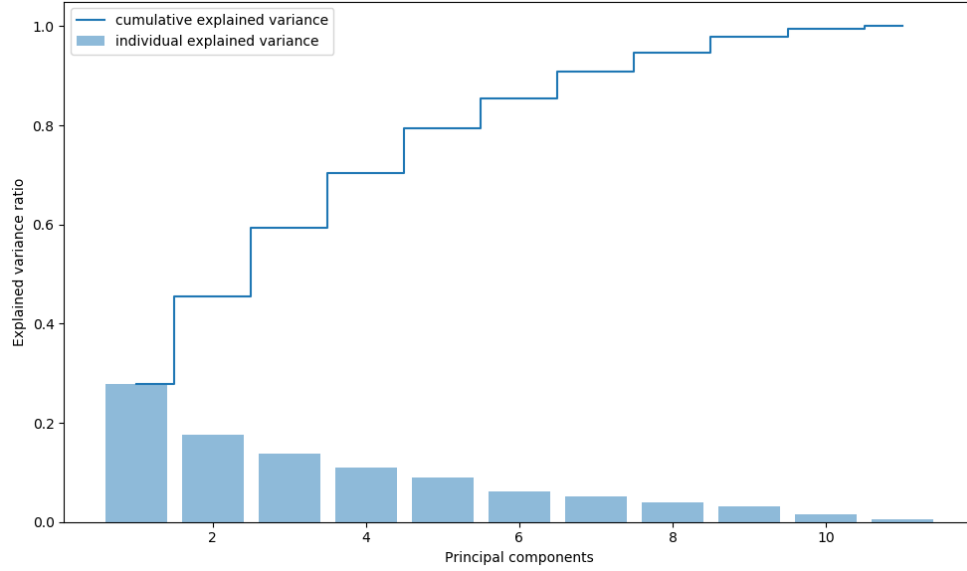


Figure 5. The individual amount of variance explained by each principal component along with the cumulative explained variance.




Figure 6. The projection of the PCA-transformed dataset into two dimensions. Each color represents a different quality score.

`max_depth = 150` and `n_estimators = 300`. Its average cross validation accuracy score was ~ 0.673 and the accuracy score was 0.652.

4.1. Outliers

In Section 3 we discussed the Elliptical Envelope and Isolation Forest algorithms for outlier detection. We concluded, based on the distributions of outliers in Figures 8 and 9, that the Isolation Forest algorithm is the superior method for our dataset. We now measure the performance of the Random Forest classifier on the dataset with the outliers detected by the Isolation Forest removed




Figure 7. Same as Figure 6 but now the data is projected into three dimensions.




Figure 8. The projection of the wines dataset onto the first three principal components with inlier data marked as blue circles and outliers as red x's. The `sklearn.covariance.EllipticalEnvelope` algorithm was used to determine the outliers. The contamination fraction was 0.04. Plot idea from [Sawtelle 2017](#).

from the training set. The document which contains the performance information of the machine learning classifiers with outliers removed is `classifier_performance_no_outliers.dat`. The best performing classifier is the Random Forest with `max_depth = 100` and `n_estimators = 300`. The accuracy score for this classifier is ~ 0.669 . We do not report the average cross-validation because the removal of the outliers likely biases the result.

REFERENCES

- | | |
|---|--|
| <p>Cortez, P., Cerdeira, A., Almeida, F., Matos, T., & Reis, J., Modeling wine preferences by data mining from physicochemical properties, In <i>Decision Support Systems</i>, Elsevier, 47(4):547-553, 2009.</p> | <p>Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., Duchesnay, E., <i>Journal of Machine Learning Research</i>. 12, 2825–2830.</p> |
|---|--|




Figure 9. The projection of the wines dataset onto the first three principal components with inlier data marked as blue circles and outliers as red x's. The `sklearn.ensemble.IsolationForest` algorithm was used to determine the outliers. The contamination fraction was 0.04. Plot idea from [Sawtelle 2017](#).

Li, Susan, Predict Employee Turnover With Python, Towards Data Science, <https://towardsdatascience.com/predict-employee-turnover-with-python-da4975588aa3>.
 Hong K, SCIKIT-LEARN : DATA COMPRESSION VIA DIMENSIONALITY REDUCTION I - PRINCIPAL COMPONENT ANALYSIS (PCA), BogoToBogo.com

Raschka, Sebastian, Principal Component Analysis in 3 Simple Steps, sebastianraschka.com
 Liu, F.T., Ting, K.M., & Zhou, Z.-H., ICDM08, 413-422, 2008.
 Sawtelle, S., <http://sdsawtelle.github.io>, Anomaly Detection in Scikit-Learn, 2017.